

HBase

RDBMS

HBASE

- Requires SQL
- Fixed schema
- Row oriented
- Not Scalable
- static
- slower retrieval of data
- It follows ACID prop.

- NO SQL
- NO fixed schema
- Column oriented
- Scalable
- dynamic
- fast retrieval
- It follows CAP (Consistency, Availability, Partition-tolerance) theorem
- Handles structured, semi-str. & unstruc. data
- Can handle sparse data.

- handles structured data

- Cannot handle sparse data

- latency is present

low latency

* Limitations of hadoop

- Performs batch processing, sequential manner
- Needed to search entire dataset.
- Processing a huge dataset results into huge data set, which is also processed sequentially

HBase.

- Open-source, horizontally scalable.
- distributed column-oriented db build on top of hadoop FS.
- Provides random real-time read/write access to data in hadoop FS.

HDFS

HBase

- for storing files

storing in HDFS & processing.

- does not support fast individual record lookups
- High latency batch processing

Provides fast lookup for larger tables. Low latency access to single row (Random access)

- Sequential access of data

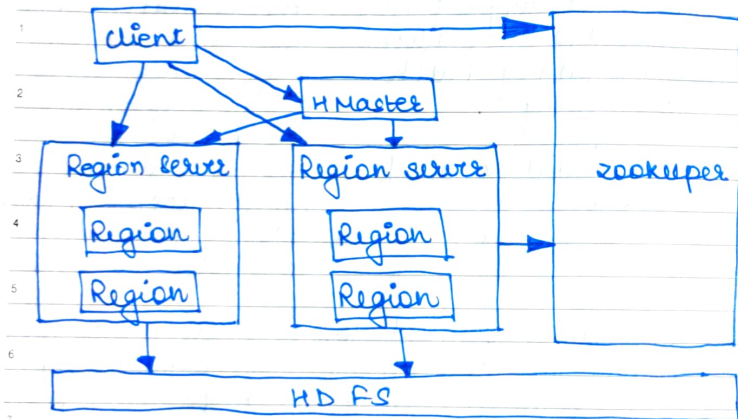
Uses hash tables & provide random access, & store data in indexed HDFS for faster lookups.

- It is column-oriented db. and table are sorted by row. The table schema only provides defines column families, which are key-value pairs.

- Table is a collection of rows
- Row is a collection of Column families
- Column family is a collection of columns
- Column is a collection of key-value pairs.

- Features of HBase
 - Linearly scalable.
 - Automatic failure support.
 - Provides consistent reads & writes.
 - Integrates with Hadoop, both as source & destn.
 - Has easy JAVA API for client
 - Provides data replication across clusters.

* Architecture



1. HMaster :-

- Master server of HBase
- Process in which regions are assigned to region server as well as DDL operations.
 - Table (create table, remove, enable/disable)
 - Column family (add column, modify column)
 - Region (move, assign)

- Manage region server instances present in cluster.
- Controlling load balancing, failover etc.
- Acts as an interface for DDL operations.

2. Zookeeper :-

- maintains server state, coordination service.
- Provides services like maintaining configuration information, naming, providing distributed synchronization, server failure notification etc.
- Clients communicate with region servers via zookeeper

3. Region Server

- HBase Tables are divided horizontally by row key range into Regions.
- Regions are the basic building elements of HBase cluster that consists of distribution of tables & are comprised of column families.
- Runs on HDFS DN.
- Regions are responsible for handling, managing, executing as well as reads & writes HBase operations on that set of region.
- Default size of region → 256 MB.

→ HMaster → Region server

- Monitoring & managing regions
- splitting regions automatically
- Handling read & writes requests.
- communicates with client directly

05

SATURDAY

FEBRUARY

- Region → building element of HBase consist of distributed tables & comprised of column families. contain multiple store, one for each column family. has memstore, n file.

★ MetaTable. (stored on Region Servers)

→ It holds the location of the Regions in HBase cluster.

→ It keeps list of all Regions in the system

→ Structure of Meta table:

1. Key: region start key, region id.
2. Values: Region Server.

→ Handled by Zookeeper.

- It is like a binary tree.

★ HBase Write Steps

① First Step: Is to write the data to the write-ahead log, while the client issues a put request:

→ To the end of WAL file, all the edits are appended which is stored on disk.

→ In case a server crashes, WAL is used to recover not-yet-persisted data

② As soon as data is written to WAL, it is stored in memstore. (256 mb written in → H Files)
in chunks 100mb

★ Region Server Components

1. WAL (Write-ahead log)

→ stores new data to permanent storage.
→ Also used for recovery.

2. Block Cache

→ It is the read cache.

FEB 2022
MTWTFSSMTWTFSSMTWTFSSMTWTFSS
• 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 • • • • •

2022

FEBRUARY

DAY: 037-328 WEEK: 06
SUNDAY

06

read

→ stores frequent data in memory & removes least read data when full.

3. MemStore

→ Write Cache, stores new data which has not been written to disk.

→ sorts before writing.

→ One MemStore per column family

4. HFiles

→ These files stores the rows as sorted key/val on disk.

★ Compaction

→ Combining of HFiles to reduce the storage & reduce the no. of disk seeks needed for read.

1. Minor: Picks smaller HFiles & combines them to bigger Hfile. It performs merge sort.

2. Major: The bigger Hfile the same column families are placed together.

→ Scheduled during low peak load timings to avoid congestion.

MAR 2022
MTWTFSSMTWTFSSMTWTFSSMTWTFSS
• 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 • • •

* 8 HBASE Architecture

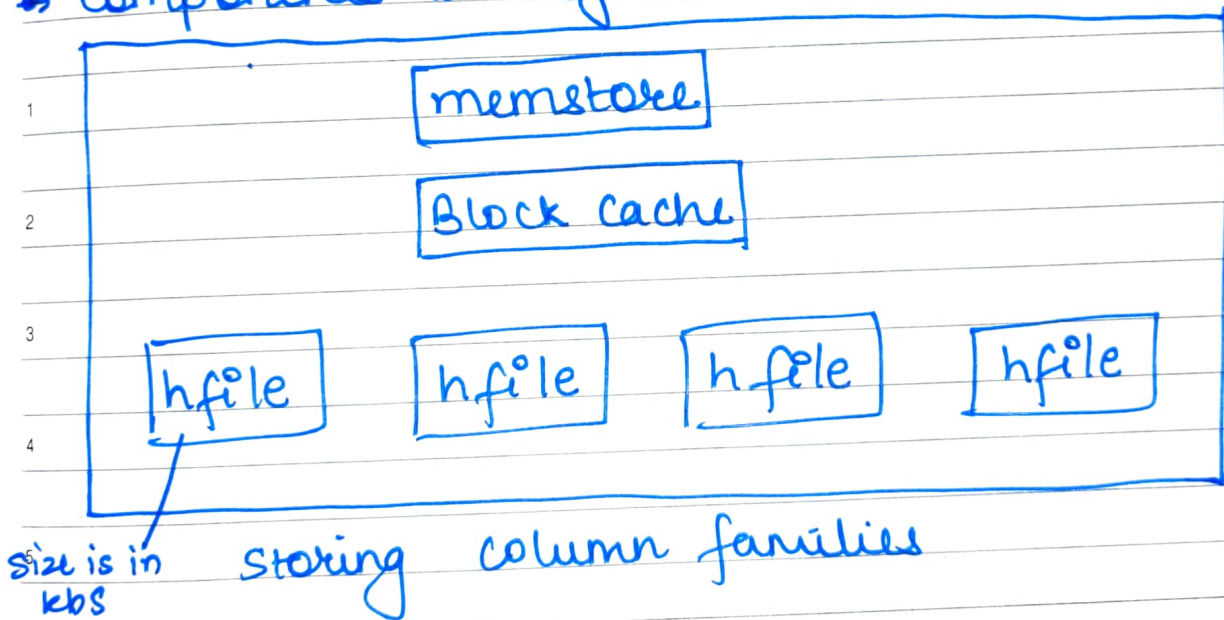
8 HBASE Architecture

9 • Zookeeper → Hmaster → Region Server
→ Regions → Column Family Store.

- Region

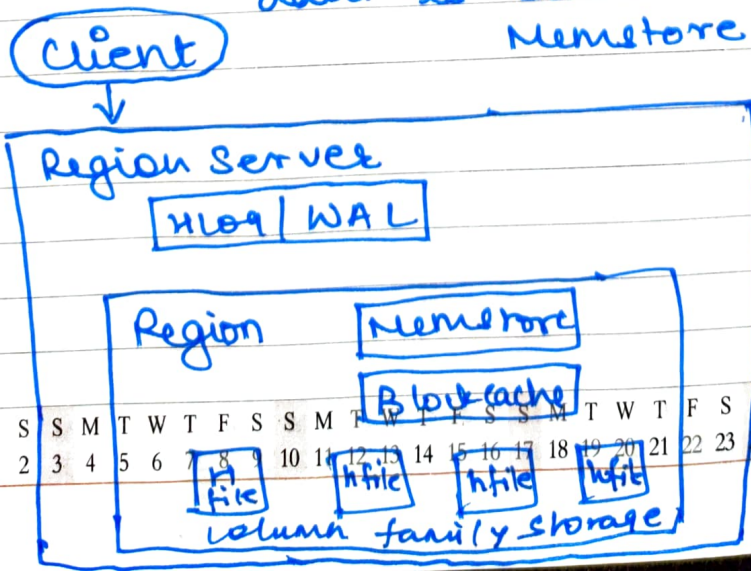
- Region
 - Actual data is splitted in regions
 - size (default) of region → 256 mb. (for optimal performance).

→ Componente in Region



Write operation

data is sent to WAL & then to Memstore → after memory is full of memstore it flushes data into Hfiles.



Before writing
data on disk it is
written in
buffer
which is memore
(assume 100 mb)

→ Read Operation

8

client → Metatable on Region Servers &

9

→ It gives location of region where table is stored → memstore, Block

10

cache, hfiles are scanned for data. → given to client.

11

12

1

2

3

