★ Kafka

- Open source stream-processing software platform which is used to handle the real-time data storage.
- Works as a broker b/w two parties i.e sender & receiver.
- Publish-subscribe messaging system which lets exchanging data b/w app^n, servers & processors as well.

- Messaging system
- Simple exchange of msg. b/w two or more persons, devices etc.
- Publish-subscribe mess. sys. allows sender to send/write the message & a receiver to read that msg.
- sender is known as producer
  Receiver is " " consumer.

- Streaming Process
- Processing of data in parallelly connected systems.
- Process allows diff. app. to limit 1^st execution of data, where one record executes w/o waiting for o/p of previous record.

- It stores streams of records in a fault-tolerant durable way.

- Apache Kafka Core API's

1. Producer API
- Allows/permits an app to publish streams of records to one or more topics.

2. Consumer API
- Allows an app^n to subscribe one or more topics & process the stream of records produced by them.

3. Streams API
- Allows an app^n to effectively transform i/p streams to o/p streams.
- It permits an app^n to act as a stream processor which consumes an i/p stream from one or more topics, & produce o/p stream to one or more o/p topics.

4. Connector API
- Executes the reusable producer & consumer APIs with existing data sys. & app^s.

- Topics
→ Refers to a particular heading or name given to some specific inter-related ideas.
- Refers to a category or a common name used to store & publish a particular stream of data.

→ Topics are similar to tables in DBS.
→ A ~~subscribes~~ producer publishes data to the topic & a consumer reads that data from the topic by subscribing it.

• Partitions
- topic is split into several parts which are known as partitions. And are seperated in an order.
- While creating topic we need to specify the no. of partitions.
- Each msg. gets stored into partitions with an incremental id known as its offset value.
- Offsets of a partition are infinite.

• Broker
- kafka cluster comprises of one or more servers known as brokers.
- It is a container that holds several topics with their multiple partitions.
- Identified with an integer id only
- Also known as Bootstrap broker as connection with any one broker means connection with entire cluster.

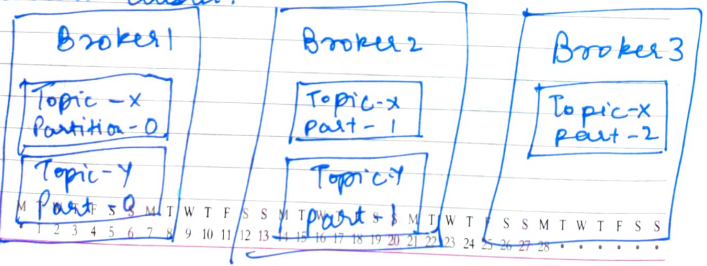

• Replication factor value should be always greater than 1.
- stores copies of Topics & partitions in other brokers
- To choose topic for client request it chooses one of the broker partition as leader & others as followers.
- If leader fails then one of the follower takes the lead.

• Zookeeper
- store info abt kafka cluster & details of consumer clients.
- Manages brokers by maintaining a list of them.
- Responsible for choosing leader.
- If any thing fails it sends msg. to kafka.
- operates with an odd no. of kafka servers
- users interact with zookeeper via brokers